

Inference latency by device (FP32, lower is better)

